

# Generative Latent Flow

Zhenya Liu

# Overview

---

1. Background

2. Generative Latent Flow

3. Experiment

## AE-Based Model

---

Consider an AE based generative model that can generate samples from the data space  $\mathcal{X}$  by the latent space  $\mathcal{Z}$ . The marginal distribution over  $\mathcal{Z}$ , denoted by  $\tilde{p}(z)$ , is unknown and depends on the Encoder  $E$ . The model also has a predefined prior distribution  $p(z)$ , and we can generate new images by sampling from the prior. Thus, in order to generate high quality samples, AE based models need to have

- a good decoder  $G$  that can output realistic images given latent variables sampled from  $\tilde{p}(z)$
- a good match between  $\tilde{p}(z)$  and  $p(z)$

## AE-Based Model

---

The first criterion is ensured by minimizing the reconstruction loss of the autoencoders. The second one can be achieved by either modifying the encoder so that  $\tilde{p}(z)$  is close to  $p(z)$ , or conversely modifying  $p(z)$  to be close to  $\tilde{p}(z)$ .

VAE adopts the first approach implicitly by using the approximated posterior  $q(z|x)$ . Training on ELBO results in minimizing  $KL[q(z)||p(z)]$ , where  $q(z) = \mathbb{E}_{x \sim p_{\text{data}}} [q(z | x)]$  (aggregated approximate posterior) is used to approximate  $\tilde{p}(z)$ .

VAEs with flow posterior have been shown to improve neither the matching of  $q(z)$  and  $p(z)$ .

## AE-Based Model

---

One problem of training  $q(z)$  to approximate  $p(z)$  is that regularizing  $q(z)$  involves one trade-off on reconstruction loss. However, we do not have such limitation when we try to learn the prior  $p(z)$  to approximate  $q(z)$ .

Also, normalizing flows can provide great flexibility to construct  $p(z)$ , and above facts motivate us to introduce the **Generative latent flow** (GLF) model.

# Architecture

---

GLF is an autoencoder model embedded with normalizing flows on the prior distribution.

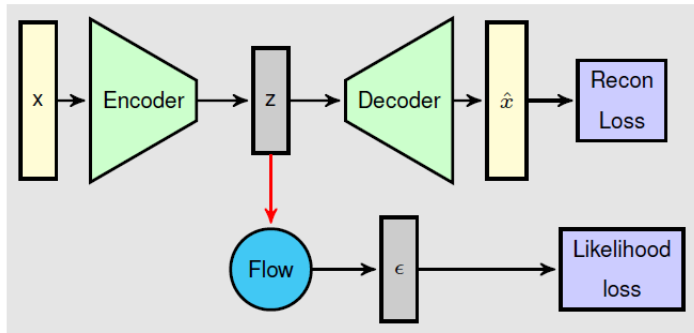


Figure: Archetecture of GLF

# Generative Latent Flow

---

Denote  $\mathcal{Z}$  be the unknown target distribution and  $\epsilon$  be the standard normal distribution. Given a bijection  $f_\theta : \mathcal{Z} \rightarrow \mathcal{E}$ ,

$$-\log(p_\theta(z)) = \mathcal{L}_{\text{NLL}}(f_\theta(z)) = - \left( \log p_\epsilon(f_\theta(z)) + \log \left| \det \left( \frac{\partial f_\theta(z)}{\partial z} \right) \right| \right)$$

# Affine Coupling Layer

---

Similar to **RealNVP**, we use affine coupling layers with alternating pattern to construct the normalizing flow  $f_\theta$ .




Figure: Affine Coupling Layers



# Generative Latent Flow

---

The loss function for GLF is composed of the reconstruction loss and NLL loss. Given encoder  $E_\eta$ , decoder  $G_\phi$ , and flow  $f_\theta : \mathcal{Z} \rightarrow \mathcal{E}$ ,

$$\mathcal{L}(\eta, \phi, \theta) = \frac{1}{N} \sum_{i=1}^N (\beta \mathcal{L}_{\text{recon}}(\mathbf{x}_i, G_\phi(E_\eta(\mathbf{x}_i))) + \mathcal{L}_{\text{NLL}}(f_\theta(\text{sg}[E_\eta(\mathbf{x}_i))]))$$

$\text{sg}[\cdot]$  is the stop gradient operation, and  $\beta$  is a hyper-parameter that controls the relative weight of two losses.

# Stop Gradient Operation

---

Note that we have  $\text{sg}[\cdot]$  term in the loss.


If we let gradients of the NLL loss back propagate into the latent variables, it can lead to degenerate of  $z$  into  $\mathbf{0}$ . That is, the encoder  $G_\eta$  will send all  $x$  to a small neighbor of 0.

This is because  $f_\theta$  has to transform the  $z$  to unit Gaussian noise, so the smaller the scale of the  $z$ 's, the more negative the log-determinant of the Jacobian becomes.


We have no control over the scale of  $z$ 's. This limitation is similar to the problem of mode collapse in GANs.

# Stop Gradient Operation

---



(a)  $p_\theta$



(b)  $p_\epsilon$

$$-\log(p_\theta(z^*)) = -\left(\log p_\epsilon(f_\theta(z^*)) + \log \left| \det \left( \frac{\partial f_\theta(z^*)}{\partial z^*} \right) \right| \right)$$

# Stop Gradient Operation

---

While the latent variables cannot become exactly 0 because of the presence of reconstruction loss in the objective, the extremely small scale of  $z$  may cause numerical issues that cause severe fluctuations.

We call our original model with stopped gradients GLF and without stopped gradients regularized GLF.

# VAEs with Flow Prior

---

The idea of GLF is closely related to VAEs with normalizing flow priors.

$$\text{ELBO}(\eta, \phi) = \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q_{\eta}(\mathbf{z}|\mathbf{x})} [\log p_{\phi, \beta}(\mathbf{x} | \mathbf{z}) + \log p(\mathbf{z}) - \log q_{\eta}(\mathbf{z} | \mathbf{x})]$$

By introducing the flow  $f_{\theta}$ ,  $p_{\theta}(\mathbf{z}) = p_{\varepsilon}(f_{\theta}(\mathbf{z})) \left| \det \left( \frac{\partial f_{\theta}(\mathbf{z})}{\partial \mathbf{z}} \right) \right|$ , ELBO becomes

$$\mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{q_{\eta}(\mathbf{z}|\mathbf{x})} \left[ \log p_{\phi, \beta}(\mathbf{x} | \mathbf{z}) + \log p_{\varepsilon}(f_{\theta}(\mathbf{z})) + \log \left| \det \left( \frac{\partial f_{\theta}(\mathbf{z})}{\partial \mathbf{z}} \right) \right| - \log q_{\eta}(\mathbf{z} | \mathbf{x}) \right]$$

## VAEs with Flow Prior

---

When the expectation over  $q_\eta(z | x)$  is estimated by sampling, this ELBO is the negative of GLF's objective (without stopping gradients) plus the entropy loss on the encoder  $q_\eta(z | x)$ .

Gaussian VAEs with flow prior does not suffer from the degeneracy of regularized GLF because of the presence of the entropy term. It is the sum of the log variances of the latent variables, and thus it encourages the encoder to output large posterior variance, preventing latent variables from collapsing to 0.

# Comparison of VAEs with flow prior on different $\beta$




Figure: [1](a) Record of FID scores on CIFAR-10 for VAEs+flow prior with different values of  $\beta$  and GLF. (b) Record of entropy losses for corresponding models. (c) Record of NLL losses for corresponding models

# Implementation of GLF

---

```
1 class AffineCoupling(nn.Module):
2     def __init__(self, input_dim, hidden_size):
3         super().__init__()
4
5         self.net = nn.Sequential(
6             nn.Linear(input_dim // 2, hidden_size),
7             nn.ReLU(),
8             nn.Linear(hidden_size, hidden_size),
9             nn.ReLU(),
10            nn.Linear(hidden_size, input_dim)
11        )
```



# Implementation of GLF

---

```
1  def forward(self, x):
2      z1, z2 = x.chunk(2, 1)
3      log_s, t = self.net(z1).chunk(2, 1)
4      s = torch.sigmoid(log_s + 2)
5      y1 = z1
6      y2 = (z2 + t) * s
7      y = torch.cat([y1, y2], dim=1)
8      logdet = torch.sum(torch.log(s).view(x.shape[0], -1), 1)
9      return y, logdet
10
11 def reverse(self, y):
12     y1, y2 = y.chunk(2, 1)
13     log_s, t = self.net(y1).chunk(2, 1)
14     s = torch.sigmoid(log_s + 2)
15     z1 = y1
16     z2 = y2 / s - t
17     x = torch.cat([z1, z2], dim=1)
18     return x
```

# Implementation of GLF

---

```
1 class Permutation(nn.Module):
2     def __init__(self, input_dim):
3         super().__init__()
4         self.input_dim = input_dim
5         self.perm = nn.Parameter(torch.randperm(input_dim),
requires_grad=False)
6
7     def forward(self, x):
8         assert x.shape[1] == self.input_dim
9         out = x[:, self.perm]
10        return out
11
12    def reverse(self, x):
13        assert x.shape[1] == self.input_dim
14        out = x[:, torch.argsort(self.perm)]
15        return out
```

# Implementation of GLF

---

```
1 class FlowNet(nn.Module):
2     def __init__(self, input_dim, hidden_size, nblocks=4):
3         super().__init__()
4
5         self.nblocks = nblocks
6
7         self.affine_layers = nn.ModuleList([AffineCoupling(input_dim,
8 hidden_size) for _ in range(nblocks)])
9
10        self.perms = nn.ModuleList([Permutation(input_dim) for _ in
11 range(nblocks - 1)])
```

# Implementation of GLF

---

```
1  def forward(self, x):
2      out = x
3      logdets = 0
4      for i in range(self.nblocks - 1):
5          out, logdet = self.affine_layers[i](out)
6          out = self.perms[i](out)
7          logdets += logdet
8      out, logdet = self.affine_layers[-1](out)
9      logdets = logdets + logdet
10     return out, logdets
11
12     def reverse(self, x):
13         out = self.affine_layers[-1].reverse(x)
14         for i in range(self.nblocks - 1):
15             out = self.perms[-1 - i].reverse(out)
16             out = self.affine_layers[-2 - i].reverse(out)
17         return out
```

# Implementation of GLF

---

```
1 class GLF(nn.Module):
2     def __init__(self, input_dim, hidden_dims, latent_dim, n_flows):
3
4         super(GLF, self).__init__()
5         self.input_dim = input_dim
6         self.hidden_dims = hidden_dims
7         self.z_dims = latent_dim
8         ## infos about using flows
9         self.n_flows = n_flows
10        ## we should create the encoder and the decoder
11        self.encoder = Encoder(input_dim, hidden_dims, latent_dim)
12        self.flows = FlowNet(latent_dim, hidden_dims, n_flows)
13        self.decoder = Decoder(latent_dim, hidden_dims, input_dim)
```

# Implementation of GLF

---

```
1  def forward(self, input):
2      # we pass the input through the encoder
3      z= self.encoder(input)
4      # we have to process the z through the flows
5      z_nograd = z.detach()
6
7      new_z, logdet = self.flows(z_nograd)
8      nll = -(log_standard_gaussian(new_z)+logdet)
9      x_reconstruct = self.decoder(z)
10
11     return x_reconstruct, nll
12
13  def sample(self, n_images):
14      # in a VAE + normalizing flows, we should start by a random
sample from N(0,1)
15      # and then we should propagate the z into the flows
16      z = torch.randn((n_images, self.z_dims), dtype = torch.float)
17      new_z = self.flows.reverse(z)
18      samples = self.decoder(new_z)
```

# Implementation of GLF

---

```
1 BATCH_SIZE = 64
2 HIDDEN_LAYERS = 200
3 Z_DIM = 40
4 N_FLOWS = 4
5
6 for i,data in enumerate(train_loader):
7     images = data[0]
8     reconstruction,nll = model(images)
9     likelihood = F.binary_cross_entropy(reconstruction, images,
reduction='sum')
10    total_loss = torch.sum(likelihood) + torch.sum(nll)
11    L = total_loss / len(images)
12    L.backward()
13    optimizer.step()
14    optimizer.zero_grad()
```

# Result

---

We have designed experiments for four types of models with VAE and flows architecture.

All encoders and decoders are two linear layers with Relu activations. All flow networks are implemented by four blocks of affine coupling layers. All training hyper-parameters are the same.




## Result

---


Model	Elbo	R_loss	KL_loss	FID
VAE	-101.55	70.64	30.85	8.03
VAE_Posterior	-102.53	78.20	24.28	7.67
VAE_Prior	-78.61	62.60	16.01	4.88
Model	Total_loss	R_loss	NLL	FID
GLF	178.63	61.97	116.57	5.24

# Reconstruction


---



(a) VAE




(b) VAE\_Posterior




(c) VAE\_PriorFlow

# Sampling


---



(d) VAE\_Posterior



(e) VAE\_Prior



(f) GLF

# Experiment

---

Now we present some results from one recurrent implementation of GLF by Ivan Fursov et al.

<https://github.com/rakhimovv/GenerativeLatentFlow>

# Experiment

---



Figure: MNIST Reconstruction after 200 epoch from GLF

We use the GLF model which was trained on CelebA dataset for 40 epochs. The dimensionality of the hidden space is 64.




Figure: Random Samples from  $\mathcal{N}(0, \mathbf{I})$



Figure: Random Samples from  $\mathcal{N}(0, \mathbf{I})$  without processing normalizing flows



Figure: Interpolation of Two Random Samples



# References

---



Xiao, Z., Yan, Q. & Amit, Y. Generative latent flow. *ArXiv Preprint ArXiv:1905.10485*. (2019)